

# Problem Set 3

Michael Fryer

2025-11-26

**Collaborators:** For this assignment, I collaborated with Florian and discussed our independent solutions.

## 1 Entropy & Divergence

```
# Calculate Information Entropy
calc_entropy <- function(l) return(-sum(l*log(l)))
```

Suppose biased coin lands heads 60% of the time.

a) What is the entropy of this coin?

```
# 60% heads, 40% tails
p <- c(0.6, 0.4)
entropy.p <- calc_entropy(p)
```

## Entropy: 0.6730117

b) Suppose another biased coin lands tails 30% of the time. Which coin has higher entropy? Justify and explain your answer.

```
# 70% heads, 30% tails
q <- c(0.7, 0.3)
entropy.q <- calc_entropy(q)
```

## Entropy: 0.6108643

The coin in part a) has a higher entropy than the coin in b). This is because the variance in the probabilities in b) is larger than in a) i.e. the more unique events there are with similar probabilities of occurring, the larger the entropy will be.

c) Suppose the first coin is the true distribution and the second coin's distribution is your approximation of the first coin. What is the KL divergence between these two distributions? Explain and justify your answer.

KL divergence is defined as:

$$D_{KL}(p, q) = \sum_i p_i \log\left(\frac{p_i}{q_i}\right)$$

```
kl <- sum(p*log(p/q))
```

## KL divergence: 0.02258242

We observe a KL divergence of 0.022. We expect this value to be small as p and q are relatively "close" together and the KL divergence is the average difference in log probabilities between the two models. This value would be much larger if instead q = {0.9, 0.1}. Importantly,  $D_{KL}(p, q) \neq D_{KL}(q, p)$  in general.

## 2 Collider Bias and Information Criteria

Return to the textbook example in §6.3.1, which explores the relationship between age, marriage and happiness. This example makes use of the simulation model `sim_happiness()`, which is part of the rethinking package:

```
d <- sim_happiness(seed=1515 , N_years=1000)
d.adults <- d[d$age > 17,] # only adults
d.adults$A <- (d.adults$age - 18) / (65 - 18)
d.adults$mid <- d.adults$married + 1
```

Compare the two models, `m6.9` and `m6.10`, using both PSIS and WAIC scores using `compare()`.

*Fit m6.9 and m6.10.*

```
# Model m6.9
m6.9 <- quap(
  alist(
    happiness ~ dnorm(mu, sigma),
    mu <- a[mid] + bA*A,
    a[mid] ~ dnorm(0, 1),
    bA ~ dnorm(0, 2),
    sigma ~ dexp(1)
  ),
  data=d.adults
)
# Model m6.10
m6.10 <- quap(
  alist(
    happiness ~ dnorm(mu, sigma),
    mu <- a + bA*A,
    a ~ dnorm(0, 1),
    bA ~ dnorm(0, 2),
    sigma ~ dexp(1)
  ),
  data=d.adults
)
```

*Score the models.*

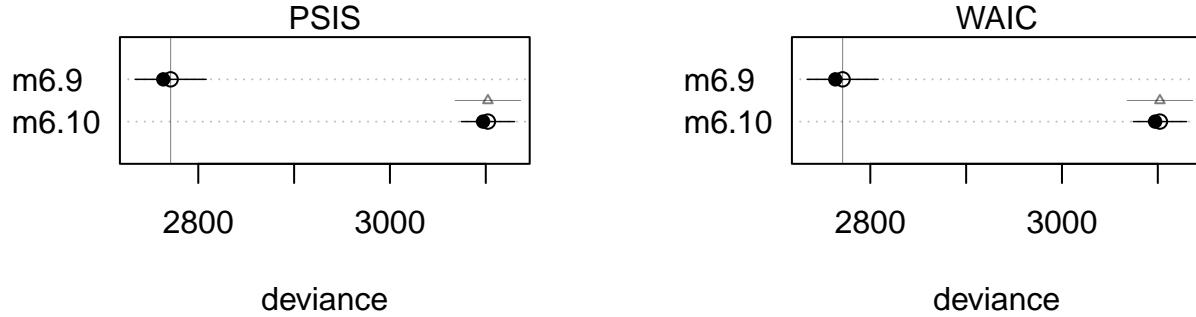
```
psis.compare <- compare(m6.9, m6.10, func=PSIS)

waic.compare <- compare(m6.9, m6.10, func=WAIC)
```

*Chart the models.*

```
##          PSIS        SE      dPSIS       dSE      pPSIS       weight
## m6.9   2771.139 37.10321  0.0000      NA 3.823583 1.000000e+00
## m6.10  3102.217 27.77867 331.0777 34.17721 2.512896 1.280557e-72

##          WAIC        SE      dWAIC       dSE      pWAIC       weight
## m6.9   2771.107 37.08184  0.0000      NA 3.807557 1.000000e+00
## m6.10  3102.197 27.76361 331.0901 34.17543 2.503099 1.272606e-72
```



- a) Which model is expected to make better predictions according to these two information criteria? Explain your answer.

*In both cross-validation approximators, a lower score is better. With that metric in mind, both WAIC and PSIS expect m6.9 to be a better prediction model. However, it should be noted that both the standard error (SE) and prediction penalty (pPSIS and pWAIC) are higher in m6.9. This means that while m6.9 is a better predictor according to WAIC and PSIS, it has more uncertainty, SE.*

- b) On the basis of the causal model, how should you interpret the parameter estimates from the model preferred by PSIS and WAIC?

*As WAIC and PSIS are cross-validation approximators, they estimate of how well a model does at **predictions**, one cannot, and should not, use them to inform any decisions about **causations**. That is, WAIC and PSIS scores have no bearing on if the model is the "correct" causal model. In fact, m6.9 is **not** the correct causal model, m6.10 is.*

### 3 Tulips and Interactions

Tulips need both water and light to bloom. But do water and light work independently, or do they interact? The `tulips` dataset in the rethinking package contains experimental data on tulip blooms under different conditions.

```
data(tulips)
d <- tulips
```

The dataset contains:

- `bed`: Bed identifier (grouping variable)
- `water`: Amount of water (1 = low, 2 = medium, 3 = high)
- `shade`: Amount of shade (1 = low, 2 = medium, 3 = high)
- `blooms`: Number of blooms produced

Before fitting any models, standardize the predictor variables and the outcome:

*Below, I do not standardize the predictor variables and the outcome. Instead, I scale the outcome to [0, 1] and center the predictor variables such that the values are -1, 0, 1. This accomplishes a similar goal as standardizing but makes it easier to reason about.*

```
# Scale blooms to [0, 1]
d$B <- (d$blooms - min(d$blooms)) / max(d$blooms)
# Center water and shade st. the values are {-1, 0, 1}
d$W <- d$water - mean(d$water)
d$S <- d$shade - mean(d$shade)
```

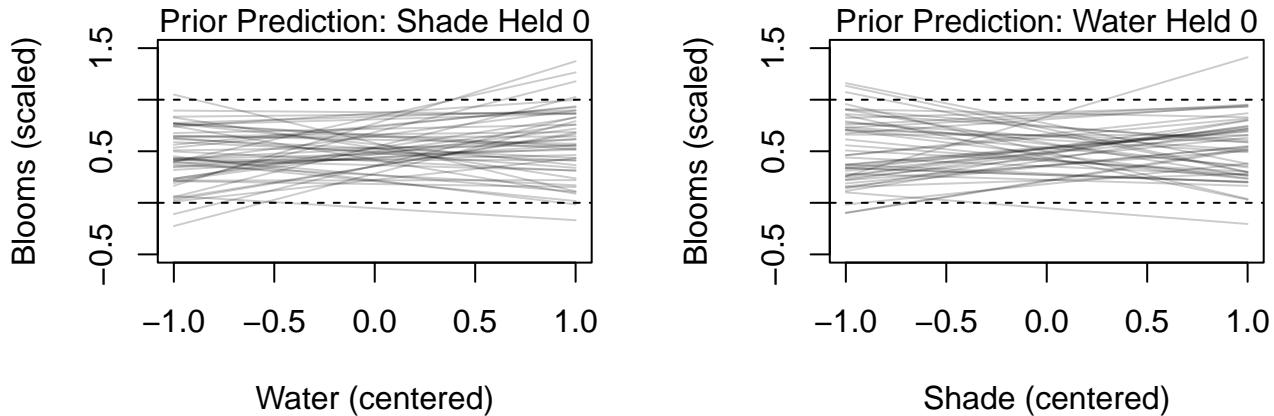
a) Fit a model (m3.1) that predicts blooms using water and shade as independent predictors (no interaction). That is, assume the effect of water on blooms is the same regardless of shade level, and vice versa.

Use prior predictive simulation to justify your priors. Write 2-3 sentences explaining your choice of priors.

**Hint:** With standardized variables, you should use relatively weak priors. A slope of 2.0 would be implausibly large for most ecological relationships.

*Since we are scaling and centering our data instead of standardizing it, our priors will not all be mean 0. We will predict our outcome prior will have  $\mu = 0.5$  as we scaled our blooms to [0, 1] and the predictor prior will have  $\mu = 0$  as we centered that data. For our standard deviations, we will use  $\sigma = 0.25$  for our outcome prior. This is because  $\pm 2\sigma$  would put blooms at impossible values for our small world i.e. < 0 or > 1. This outcome prior is arguably still loose and that is justified by the fact that we have a small dataset  $n = 27$ . For the same reasons, we will use  $\sigma = 0.25$  for our predictor variables.*

```
N <- 50
a <- rnorm(N, 0.5, 0.25)
bS <- rnorm(N, 0, 0.25)
bW <- rnorm(N, 0, 0.25)
```



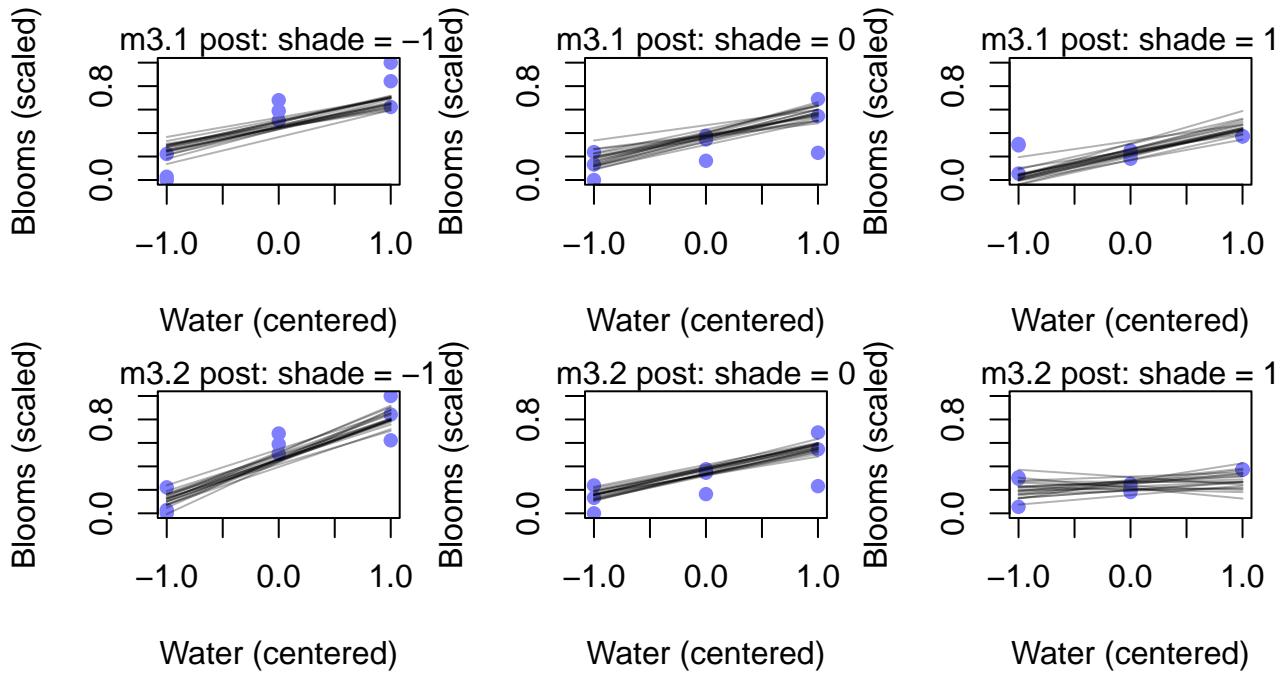
```
m3.1 <- quap(
  alist(
    B ~ dnorm(mu, sigma),
    mu <- a + bW*W + bS*S,
    a ~ dnorm(0.5, 0.25),
    bW ~ dnorm(0, 0.25),
    bS ~ dnorm(0, 0.25),
    sigma ~ dexp(1)
  ),
  data=d
)
```

b) Now fit a model (`m3.2`) that includes an interaction between water and shade. This allows the effect of water on blooms to depend on the level of shade (and vice versa).

**Hint:** In the model formula, you can specify the interaction as  $W * S$ , which is shorthand for  $W + S + W * S$ . Or you can write out all three terms explicitly.

```
m3.2 <- quap(
  alist(
    B ~ dnorm(mu, sigma),
    mu ~ a + bW*W + bS*S + bWS*W*S,
    a ~ dnorm(0.5, 0.25),
    bW ~ dnorm(0, 0.25),
    bS ~ dnorm(0, 0.25),
    bWS ~ dnorm(0, 0.25),
    sigma ~ dexp(1)
  ),
  data=d
)
```

Plot the posterior predictions for both models (`m3.1` and `m3.2`) side by side. For each model, create a plot showing how predicted blooms change with water at different levels of shade (e.g., plot three lines for shade = -1, 0, 1).



c) Use both PSIS and WAIC to compare your two models. Which model is expected to make better predictions? Is the difference substantial?

Create the comparison plot using `plot(compare(m3.1, m3.2))` and interpret what you see. Write 3-4 sentences explaining:

- Which model performs better
- How large the difference is relative to the standard error
- What this tells you about whether water and shade interact in affecting tulip blooms

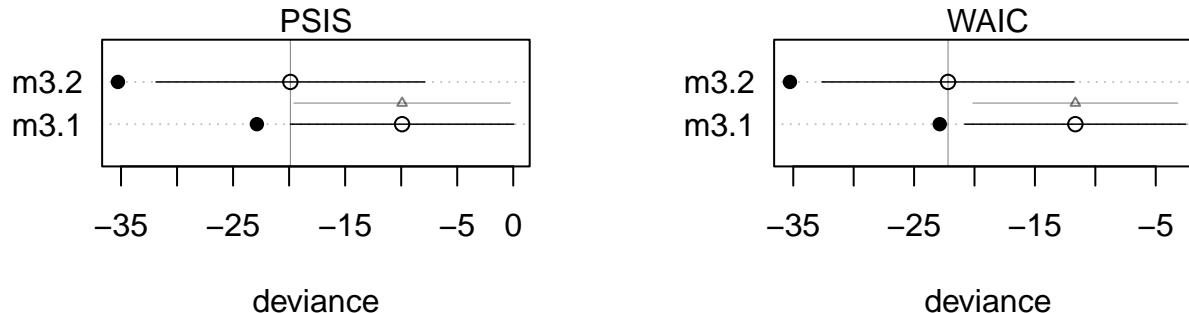
*Score the models.*

```
psis.compare <- compare(m3.1, m3.2, func=PSIS)
```

```
waic.compare <- compare(m3.1, m3.2, func=WAIC)
```

*Chart the models.*

```
##          PSIS           SE      dPSIS       dSE      pPSIS      weight
## m3.2 -19.888147 11.98897 0.000000      NA 7.693312 0.993208772
## m3.1  -9.917529  9.97304 9.970618 9.657799 6.479925 0.006791228
##          WAIC           SE      dWAIC       dSE      pWAIC      weight
## m3.2 -22.19026 10.438988 0.000000      NA 6.542255 0.994873897
## m3.1 -11.65372  9.145693 10.53654 8.481705 5.611829 0.005126103
```

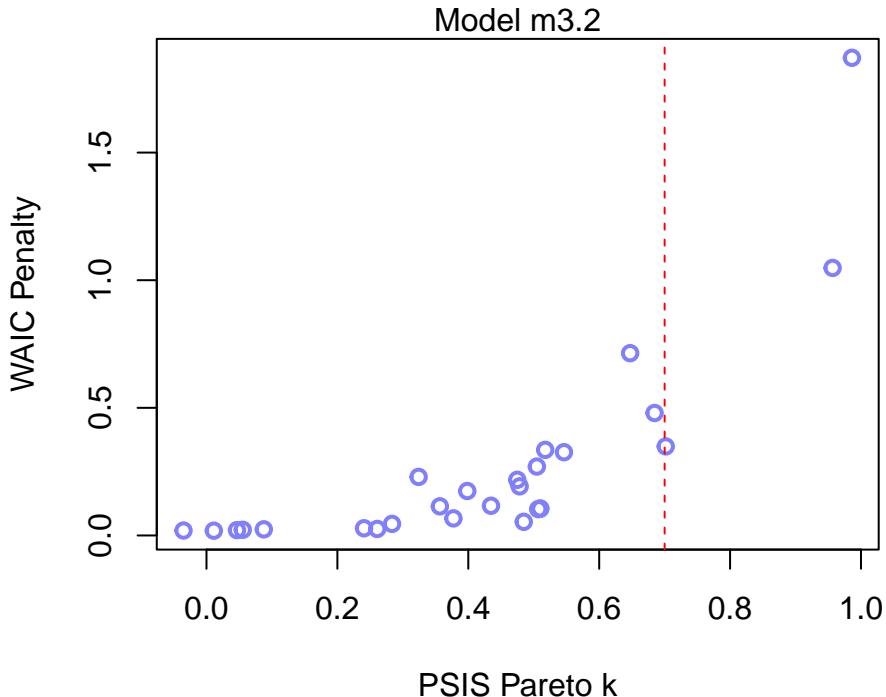


```
## 80% Confidence Interval
## PSIS:  -2.391365 22.3326
## WAIC:  -0.3200416 21.39312
```

When comparing our two models using PSIS and WAIC we see similar results. With each method, m3.2 is expected to perform better (lower values). When comparing the difference, **dPSIS** and **dWAIC**, we observe that it is about the same size as the standard error **SE** which tells us that these models are potentially not substantially different. When calculating a generous 80% confidence interval (see above) we observe that the whole 80% interval is not above 0 meaning that these two models potentially overlap. This gets more egregious with more standard CI such as 90%, 95%, and 99%. With that in mind, I would not conclude that these two models are significantly different.

d) Examine the pointwise PSIS diagnostics for your preferred model. Are there any problematic observations (high Pareto  $k$  values)? If so, which observations are influential and why might they be outliers?

**Hint:** Use `PSIS(model, pointwise=TRUE)` and examine the  $k$  values.



```
## Max Pareto k: 0.986
## k > 0.7: 3 (11.1%)
```

When there are Pareto  $k$  values above 0.5, the Pareto distribution has infinite variance. Still, we will use a cutoff of 0.7 as theory and simulation suggests values  $k < 0.7$  are still fine to use. That being said, we observe 3 such values with  $k > 0.7$  which accounts for 11.1% of our data. Additionally, one such value is nearly double the theoretical cutoff where the variance becomes infinite (0.5). This is very strong evidence towards the presence of outliers. Possible explanations for such outliers are plants that failed to produce blooms despite getting adequate sun and water.

e) Based on your analysis, write 2-3 sentences answering: Do water and light work independently to affect tulip blooms, or do they interact? What does this mean biologically?

*It is very likely that light (shade) and water interact to affect tulip blooms. When varying the levels of shade in m3.2 we observe a better fit than in the model we did not, m3.1. This interaction can be explained biologically by the fact that plants need light and nutrients to grow i.e. plants that have more access to light, need more water to support photosynthesis, and vice versa.*

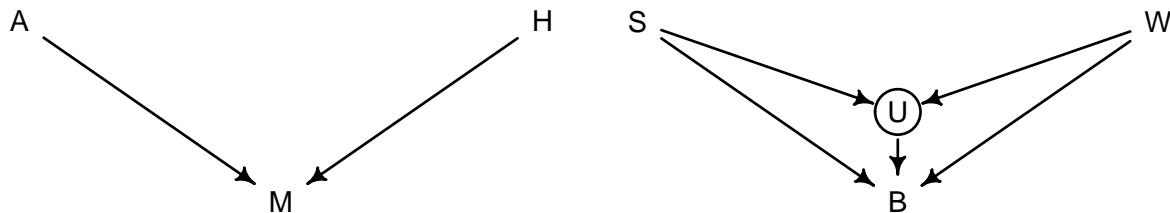
## 4 Reflection

Look back at Problems 2 and 3. Both involve interactions, but of different kinds:

- In Problem 2, you examined collider bias, where conditioning on a variable (marriage status) creates a spurious association between age and happiness
- In Problem 3, you examined multiplicative interactions, where the effect of one variable (water) depends on the level of another (shade)

Write 3-4 sentences explaining: What is the difference between these two types of “interactions”? How are they similar and how are they different from a causal modeling perspective?

*These two interactions are both a result of a collider, however the statistical relationships are different. In problem 2, the outcome has a positive relationship with both predictor variables. However, in problem 3, the outcome has a positive relationship with one predictor variable, and a negative relationship with another. Additionally, the impact of one predictor influences the affect of another resulting in a multiplicative interaction. Pictured below is a graphical representation of what I believe the causal models to look like.*



## 5 AI Declaration

AI was not used for this assignment although, I probably should have because I had a bug in my models where I had `sigma ~ exp(1)` written instead of `sigma ~ dexp(1)` so the predictions made no sense and it took me forever to find it.

## 6 Appendix

```

# Plots a standardized predictor x
plot_prior_prediction <- function(x, ylim, xlab, ylab, title="Prior Prediction") {
  plot(NULL, xlim=range(x), ylim=ylim, xlab=xlab, ylab=ylab)
  mtext(title)
  # Graph all the lines
  for (i in 1:N) {
    curve(a[i] + b[i]*x,
      from=min(x), to=max(x),
      col=col.alpha("black", 0.2), add=TRUE
    )
  }
  # Bounds of impossibilities
  abline(h=0, lty=2)
  abline(h=1, lty=2)
}

# Plot the posterior predictions for blooms
plot_post_prediction_blooms <- function(m, m.name) {
  for (s in -1:1) {
    i <- which(d$S==s)
    plot(d$W[i], d$B[i],
      xlim=c(-1, 1), ylim=c(0, 1),
      xlab="Water (centered)", ylab="Blooms (scaled)",
      pch=16, col=rangi2
    )
    mtext(sprintf("%s post: shade = %d", m.name, s))
    post <- link(m, data=data.frame(S=s, W=-1:1))
    for (j in 1:20) lines(-1:1, post[j,], col=col.alpha("black", 0.3))
  }
}

# Calculate a confidence interval
calc.interval <- function(mu, sigma, z=2.575) { # Default of 99% z-score
  return(mu + c(-1,1) * sigma * z)
}

# Plot psis k-values
psis <- PSIS(m3.2, pointwise=TRUE)
waic <- WAIC(m3.2, pointwise=TRUE)
plot(psis$k, waic$penalty,
  xlab="PSIS Pareto k", ylab="WAIC Penalty",
  col=rangi2, lwd=2
)
mtext("Model m3.2")
abline(v=0.7, col="red", lty=2)

# Plot the supposed DAGs for problems 2 and 3
p2.dag <- dagitty("dag {
  A -> M
  H -> M
}")
coordinates(p2.dag) <- list(x=c(A=-0.5, H=0.5, M=0), y=c(A=0, H=0, M=0.5))

```

```

drawdag(p2.dag, xlim=c(-1, 1), ylim=c(-1, 1))
p3.dag <- dagitty("dag {
  S -> B
  W -> B
  S -> U [unobserved]
  W -> U
  U -> B
}")
coordinates(p3.dag) <- list(x=c(S=-0.5, W=0.5, U=0, B=0), y=c(S=0, W=0, U=0.25, B=0.5))
drawdag(p3.dag, xlim=c(-1, 1), ylim=c(-1, 1), shapes=c(U="c")))

```